

Using the Dremio Connector for Tableau



October 20th, 2020

GREG PALMER

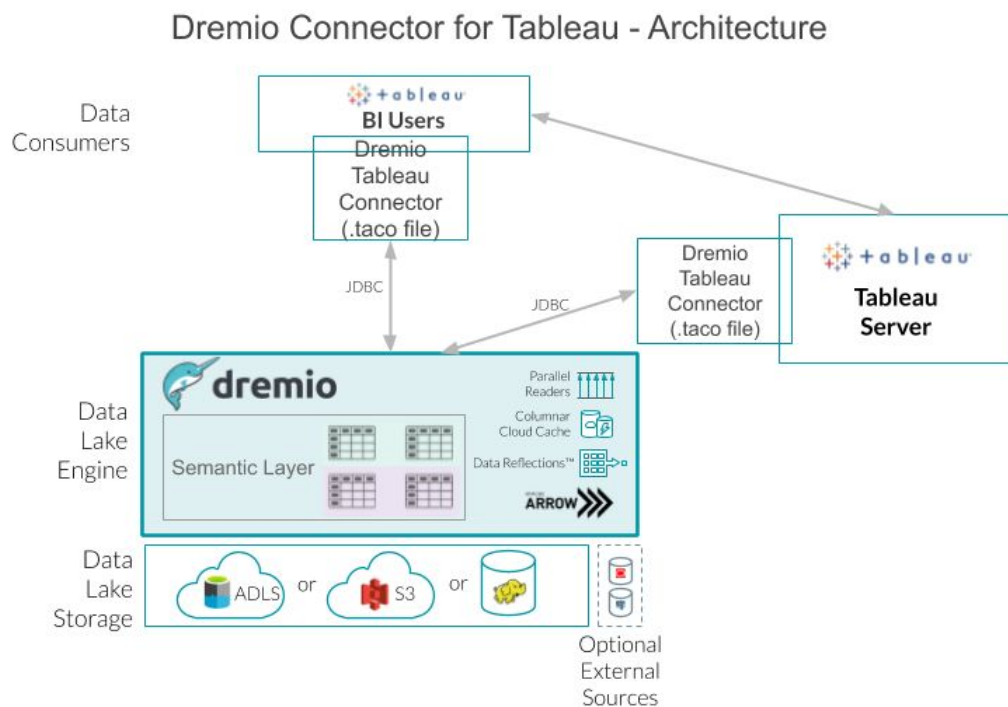
In the past, Tableau users utilized the Dremio ODBC driver to connect to the Dremio Data Lake Engine from their Tableau workbooks using the “Other Database” option. Now Dremio offers a customized connector for Tableau that is available from the [Tableau Extension Gallery](#). This new connector was created using the Tableau Connector SDK and offers better performance and more features over the basic driver, including:

- Better live query support. Allows a customized SQL dialect to generate queries that are optimized for Dremio.
- Simpler connection experience. Users don’t need to enter obscure JDBC URL strings or create a DSN or configure odbc.ini files. The connector provides a simple Dremio connection dialog.
- Runs in Tableau Desktop and Tableau Server. No configuration is required once you install the connector.
- Enables the use of Tableau’s “Impersonate using embedded password” feature combined with Dremio’s “Inbound Impersonation” feature.

To use the Dremio Tableau Connector, it must be deployed on each users’ Tableau Desktop computer and on the Tableau Server computer. When a user creates a new Tableau workbook using Tableau Desktop, they can specify the

Dremio Connector instead of using the “Other Database” driver. When a user publishes a workbook to the Tableau Server, that same Dremio Connector will be used to connect to Dremio when the workbook is opened.

The Dremio Connector for Tableau architecture is illustrated in the following diagram.



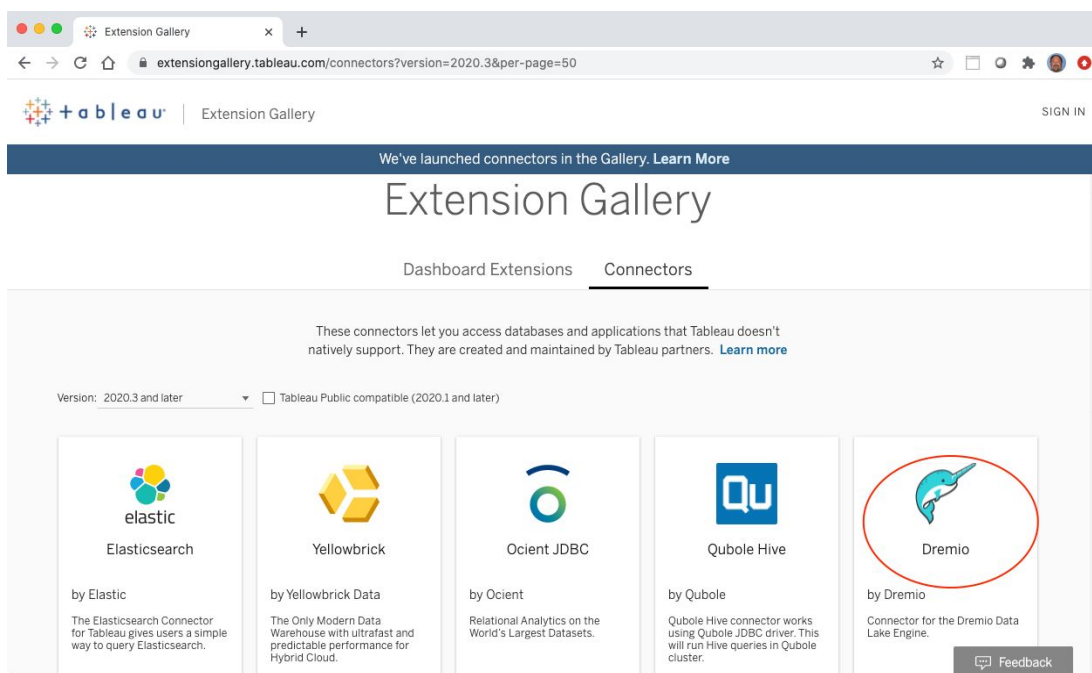
The steps required to install and use the new Dremio Connector for Tableau include the following:

1. Install the Dremio Connector on Tableau Desktop
2. Install the Dremio Connector on Tableau Server
3. Configure the Dremio “Inbound Impersonation” capability and .tds file generation
4. Create a Tableau data source using the Dremio Connector
5. Create and publish a Tableau workbook using Tableau’s “Impersonate using embedded password” configuration

6. Run a Tableau workbook in Tableau Server as a different user using Dremio's "Inbound Impersonation" configuration

1. Install the Connector on Tableau Desktop

Download the Dremio Tableau Connector from the Tableau Extension Gallery at <https://extensiongallery.tableau.com/connectors>. Click on the Dremio connector and then click on the blue Download button.



Sign in with your Tableau username and password (or create an account) and download the `dremio.taco` file to your Web browsers download directory.

Move the `dremio.taco` file to your Tableau Desktop connectors directory. If the directory does not exist, create it.

On Windows move the file to:

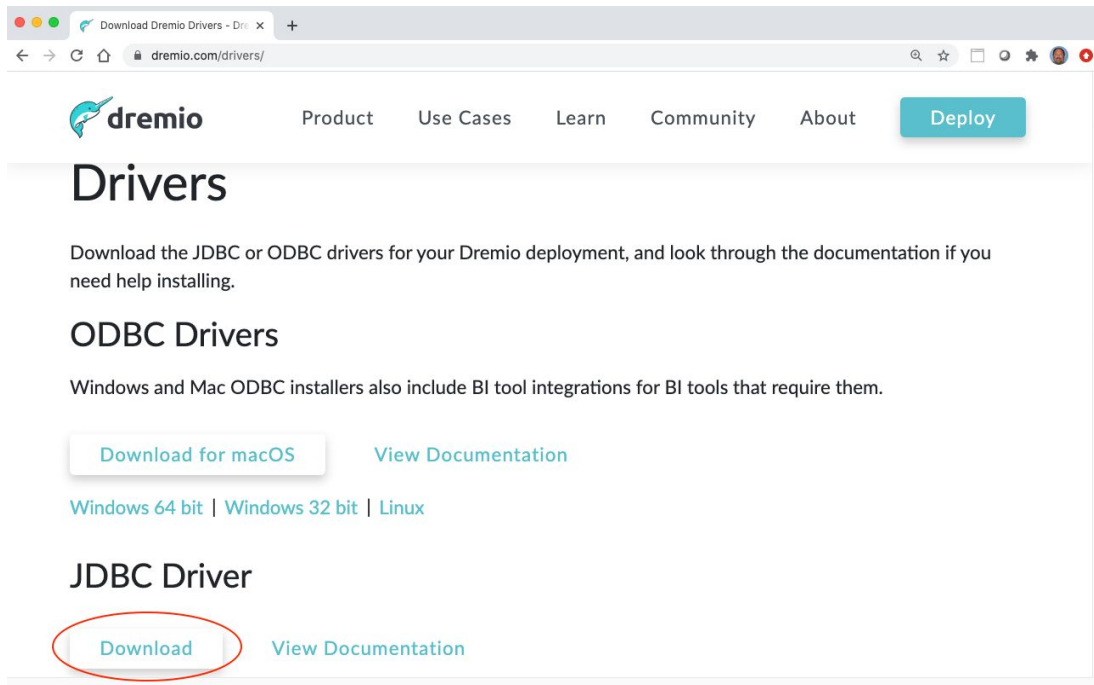
```
\Users\[Windows User]\Documents\My Tableau Repository\Connectors
```

On macOS move the file to:

dremio.com | +1.650.383.6805 | contact@dremio.com

```
/Users/[user]/Documents/My Tableau Repository/Connectors
```

Now download the Dremio JDBC driver from the Dremio website at <https://www.dremio.com/drivers>.



Click on the Download link and download the `dremio-jdbc-driver-{version}.jar` file to your Web browser's download directory.

Move the `.jar` file to your Tableau Desktop drivers directory. If the directory does not exist, create it.

On Windows move the file to:

```
C:\Program Files\Tableau\Drivers
```

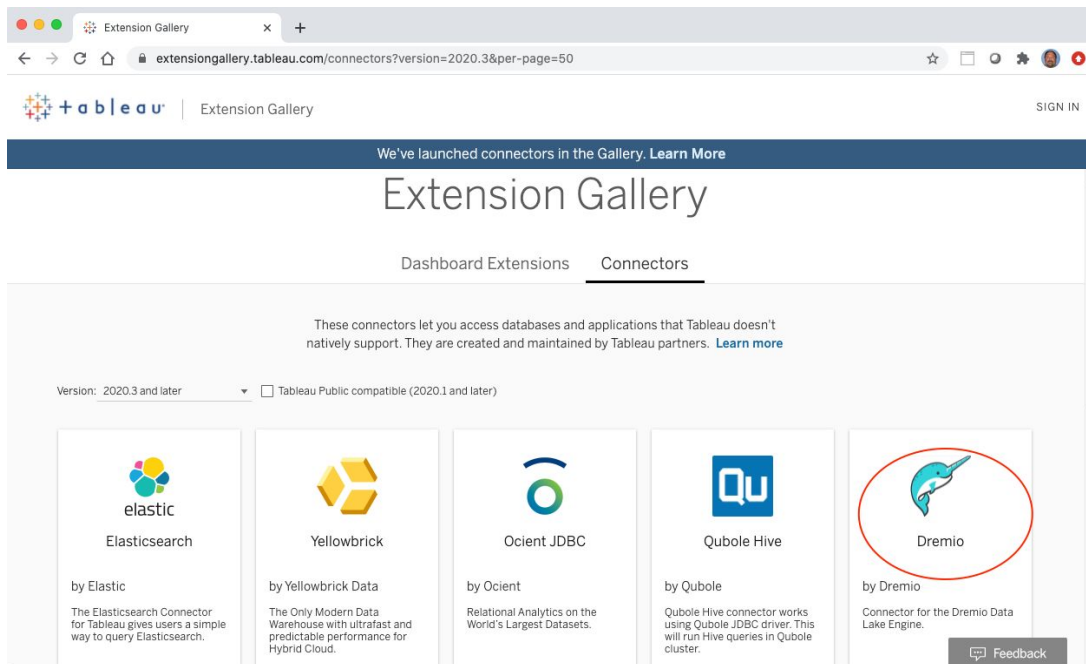
On macOS move the file to:

```
/Users/[user]/Library/Tableau/Drivers
```

Finally, restart your Tableau Desktop application.

2. Install the Connector on Tableau Server

Download the Dremio Tableau Connector from the Tableau Extension Gallery at <https://extensiongallery.tableau.com/connectors>. Click on the Dremio connector and then click on the blue **Download** button.



Sign in with your Tableau username and password (or create an account) and download the `dremio.taco` file to your Web browsers download directory.

Move the `dremio.taco` file to your Tableau Server connectors directory. If the directory does not exist, create it.

Move the `.taco` file to:

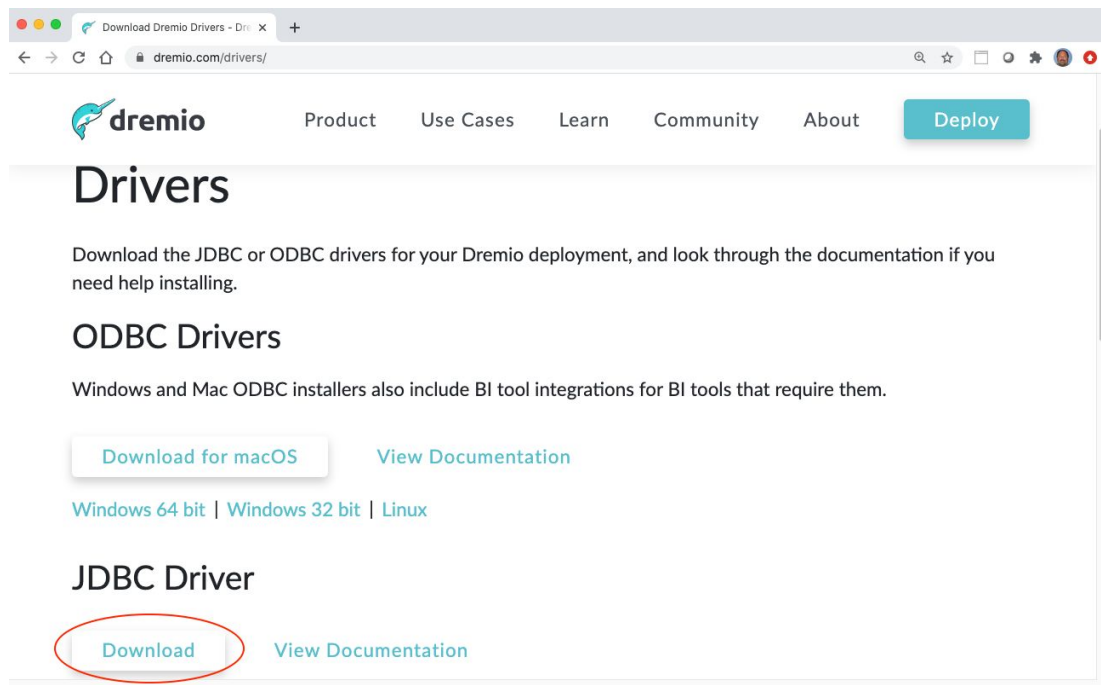
```
[Tableau_Server_Install_Dir]\data\tabsvc\vizqlserver\Connectors
```

In most cases, the Tableau Server is installed in the Program Data directory, so you would copy the `.taco` file to:

dremio.com | +1.650.383.6805 | contact@dremio.com

```
C:\ProgramData\Tableau\Tableau Server\  
data\tabsvc\vizqlserver\Connectors
```

Now download the Dremio JDBC driver from the Dremio website at <https://www.dremio.com/drivers>.



Click on the Download link and download the `dremio-jdbc-driver-{version}.jar` file to your Web browser's download directory.

Move the `.jar` file to your Tableau Server drivers directory. If the directory does not exist, create it. Move the file to:

```
C:\Program Files\Tableau\Drivers
```

Note: If you are running Tableau Server 2030.3+ that uses the Java 11 runtime for JDBC drivers, you must add an environment variable to enable Dremio to function properly. Follow these steps:

- Login as the user that runs the Tableau Server
- From Windows Server desktop, go to Control Panel > System Security > System > Advanced system settings
- Click the **Environment Variables** button in the window that pops up
- Add a new variable under User variables for <user that runs Tableau>
- Create the new variable name:
`JAVA_TOOL_OPTIONS`
- Set the new variable value to:
`-Dcdjd.io.netty.tryReflectionSetAccessible=true`

Finally, restart your Tableau Server Windows system service.

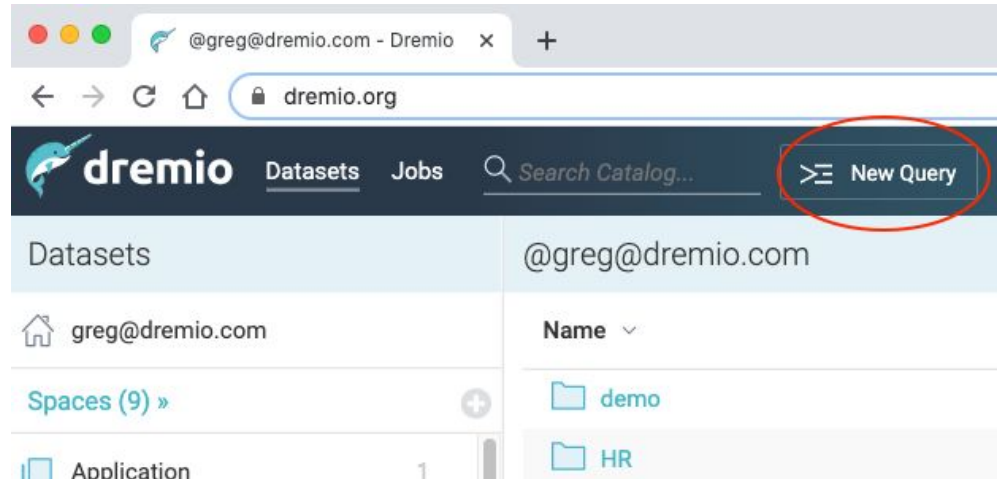
3. Configure the Dremio “Inbound Impersonation” capability and .tds file generation

Inbound Impersonation

Dremio supports the concept of inbound impersonation, where the end-user connects to Dremio using a proxy user. A proxy user is a user that is allowed to "connect on behalf of another user". In this case, Dremio runs a query as the real user that submitted the request (and verifies any access permissions needed).

To configure Dremio to enable proxy users that submit queries as another user, a Dremio Administrator must create an inbound policy via a Dremio support key named `exec.impersonation.inbound_policies`. Follow these steps:

- Log into the Dremio UI as an Administrator user
- Launch a New Query session by clicking on the **New Query** button:



- In the query editor window, type in an `ALTER SYSTEM` command with the following format:


```
ALTER SYSTEM SET
"exec.impersonation.inbound_policies"=' [
{
  proxy_principals:{
    Users:["<proxy-user-1>"]
  },
  target_principals: {
    users:["<real-user-1>",
           "<real-user-2>"]
  },
  proxy_principals:{
    Users:["<proxy-user-2>"]
  },
  target_principals: {
    users:["<real-user-2>",
           "<real-user-3>",
           "<real-user-4>"]
  }
}
```



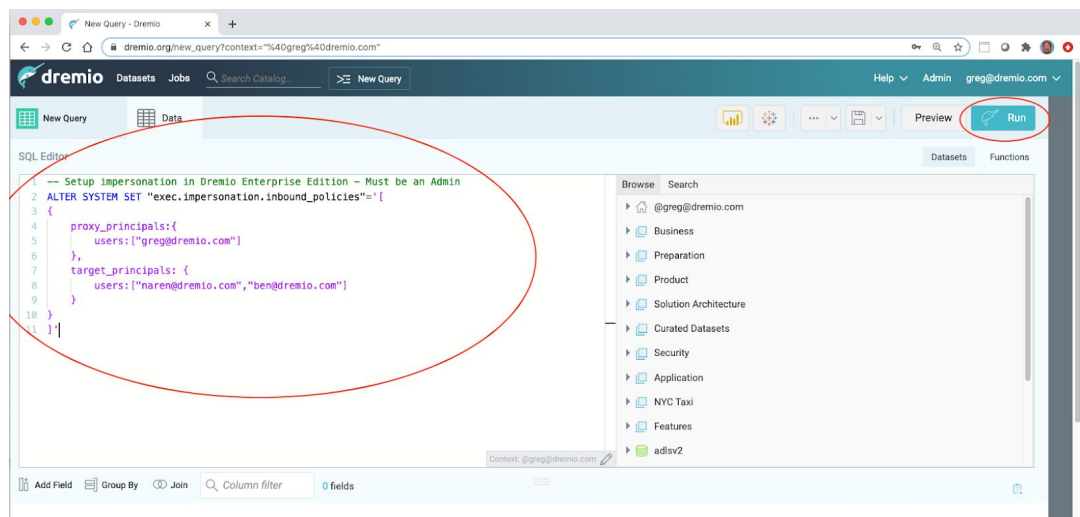
```
}
]'
```

Where the proxy user `proxy-user-1` can submit Dremio queries for the real users `real-user-1` and `real-user-2`. And where the proxy user `proxy-user-2` can submit Dremio queries for the real users `real-user2`, `real-user3` and `real-user-4`.

- Click on the **Run** button to run the `ALTER SYSTEM` command:



A real example might look like this:



- After running the `ALTER SYSTEM` command, you can verify that it was effective by running the command:

```
-- Verify the ALTER SYSTEM command
```

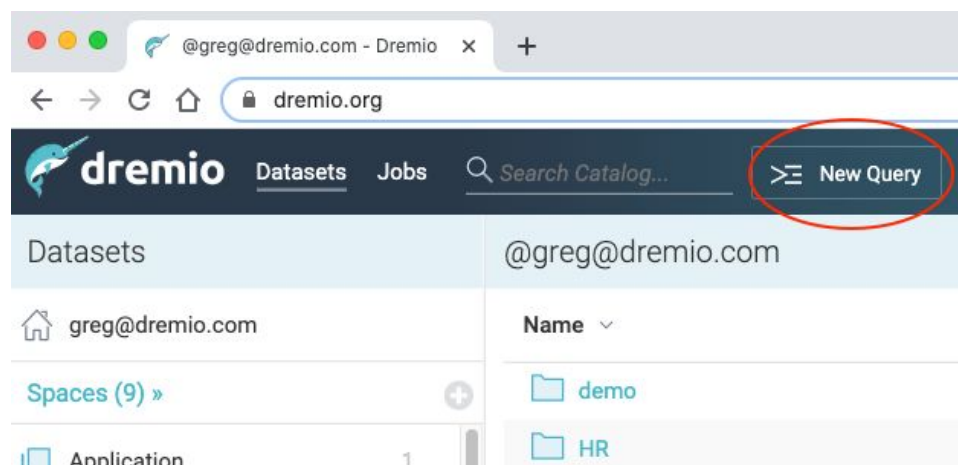
```
SELECT name, string_val
FROM sys.options
WHERE name = 'exec.impersonation.inbound_policies'
```

Tableau .tds File Generation

The Dremio UI allows users to quickly launch a Tableau Desktop session by generating a .tds file that contains the connection information to the Dremio server. This .tds file gets downloaded to the user's web browser download directory and then can be "launched" and create a new Tableau workbook.

By default, Dremio configures the generated .tds file with a generic ODBC connection specification (`connection class="genericodbc"`). To enable the Dremio Connector for Tableau (`connection class="dremio"`), a Dremio Administrator must configure a "Support Key" named `export.tableau.export-type`. Follow these steps:


- Log into the Dremio UI as an Administrator user
- Launch a New Query session by clicking on the **New Query** button:



- In the query editor window, type in an `ALTER SYSTEM` command with the following format:

```
-- Configure Dremio to use the Dremio Connector
-- instead of generic ODBC in .tds files
ALTER SYSTEM
    SET "export.tableau.export-type" = 'NATIVE'
```

- Click on the **Run** button to run the `ALTER SYSTEM` command:



- After running the `ALTER SYSTEM` command, you can verify that it was effective by running the command:

```
-- Verify the ALTER SYSTEM command
SELECT name, string_val
FROM    sys.options
WHERE   name = 'export.tableau.export-type'
```


Finally, restart your Dremio server by launching an SSH session into the Dremio Coordinator Node server and running the command:

```
sudo systemctl restart dremio
```

4. Create a Tableau data source using the Dremio Connector


Now that the Dremio Connector for Tableau has been installed on the Tableau Desktop computer and the Tableau Server computer, a Tableau data source can be created that uses the new connector. Let's start by launching Tableau from the Dremio UI.

Log into the Dremio Web UI as either a non-admin user or an admin user. If you don't have a Dremio space named `Preparation`, create one by clicking on the Add Space button or the plus sign icon to the right of the Spaces.




In the **Add Space** screen, enter the name `Preparation` and click the **Save** button.


Now create a new Physical Dataset from the Dremio Samples dataset. Click on the **Add Sample Source** button.



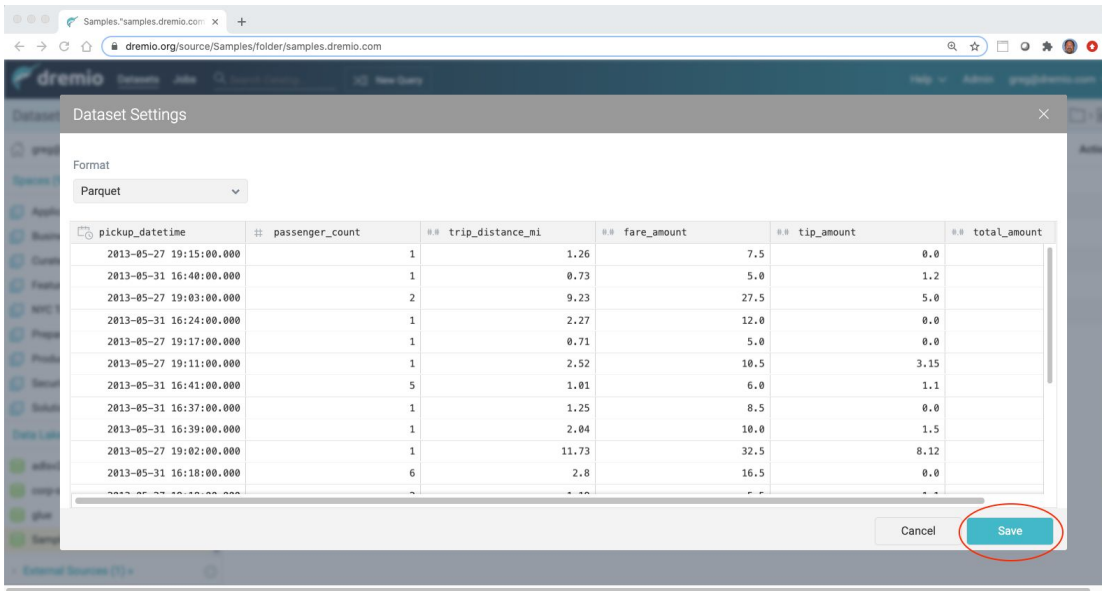
Then click on the sample folder that is generated.



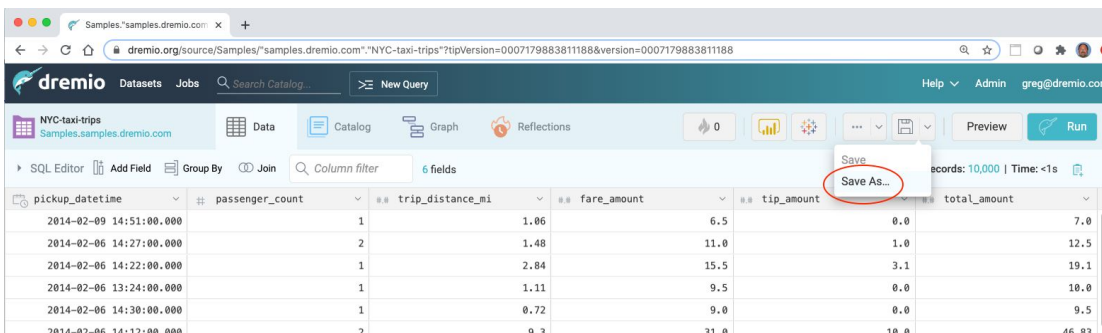
Inside the folder, click on the **Format Folder** icon in the **Actions** for the `NYC-Taxi-Trips` dataset.



The sample data is stored using the Parquet file format. In the **Dataset Settings** screen click the **Save** button to save this default formatting.



Then click on the **Save As** button to save this as a new VDS.



In the **Save Dataset As** screen, select the **Preparation** space and enter the name `sample_nyc_trips` and click on the **Save** button.

Name

sample_nyc_trips

Location

@greg

Preparation

Cancel Save

Speed up query performance against this dataset by creating a new Dremio Aggregation Reflection. Click on the **Reflections** button and enable the default Aggregation Reflection. Then click the **Save** button to begin generating this new reflection.

dremio Datasets Jobs Search Catalog... New Query Help Admin greg

sample_nyc_trips Preparation Data Catalog Graph Reflections 0 Preview Run

Reflections Switch to Advanced

Raw Reflections

Aggregation Reflections

"sample_nyc_trips" dataset fields:

Search fields...

pickup_datetime

passenger_count

trip_distance_mi

fare_amount

tip_amount

total_amount

Dimensions

passenger_count

pickup_datetime

Clear All

Measures

trip_distance_mi

total_amount

tip_amount

fare_amount

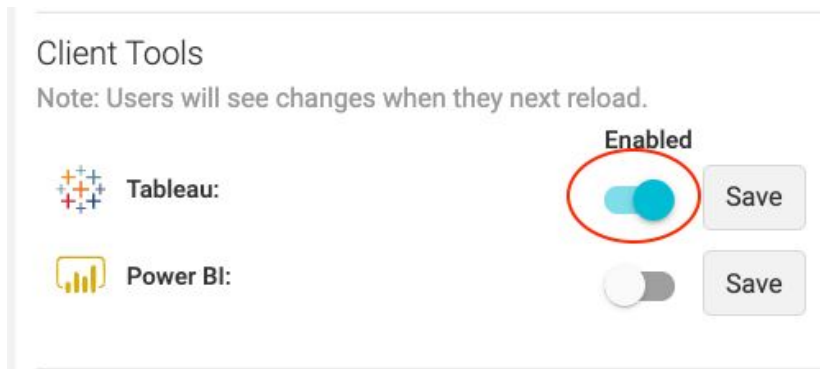
Clear All

Add a Dimension Add a Measure

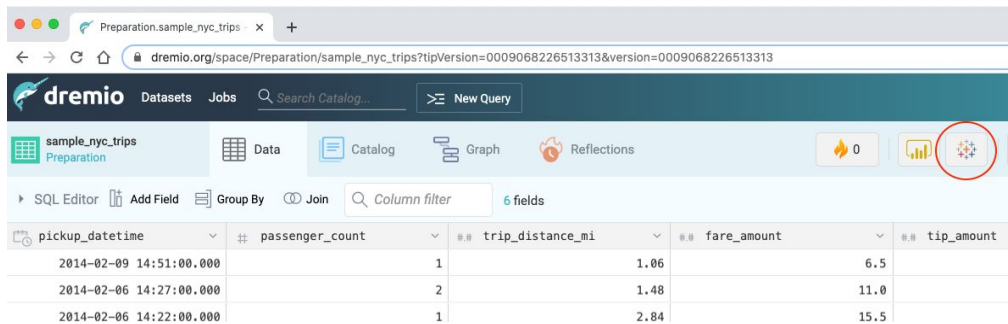
Revert Save

Now launch Tableau Desktop using the .tds file generated by Dremio. In the sample_nyc_trips dataset preview screen, click on the **Tableau** button.

The next step requires the use of the **Tableau** launch button. If you do not see a **Tableau** button, you can enable it as a Dremio Administrator user by clicking on the **Admin** link at the top right of the page and then click on the **Support** link on the left side of the page. In the **Client Tools** section, click on the **Tableau** toggle button and then click the **Save** button.



Now launch Tableau Desktop using the .tds file generated by Dremio. In the `sample_nyc_trips` dataset preview screen, click on the **Tableau** button.

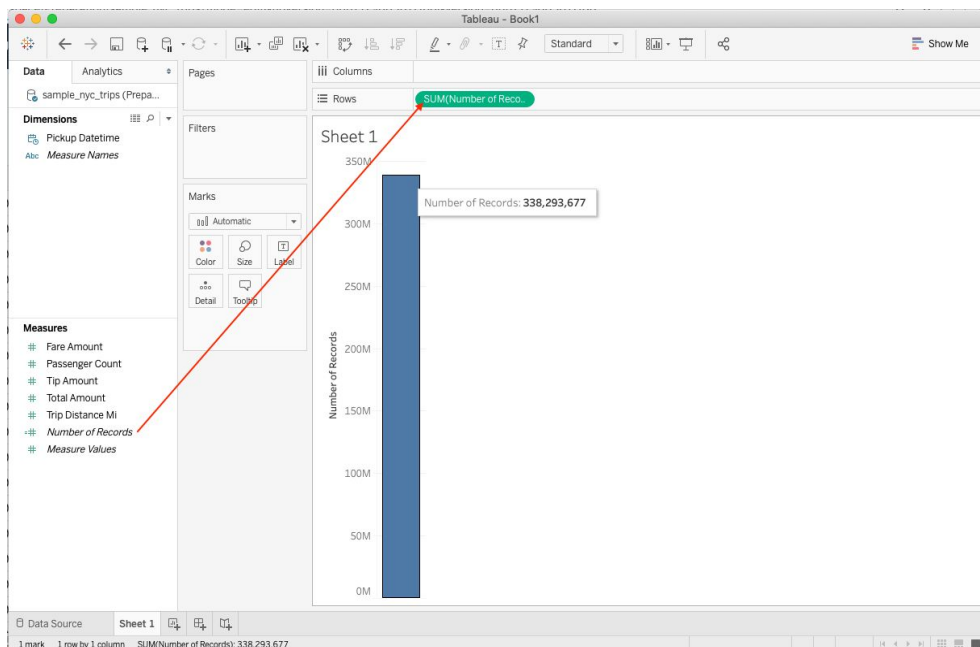


This will download a .tds file in your Web browser's download directory. At the bottom of your Web browser page, click on the downloaded file and **Open** it.

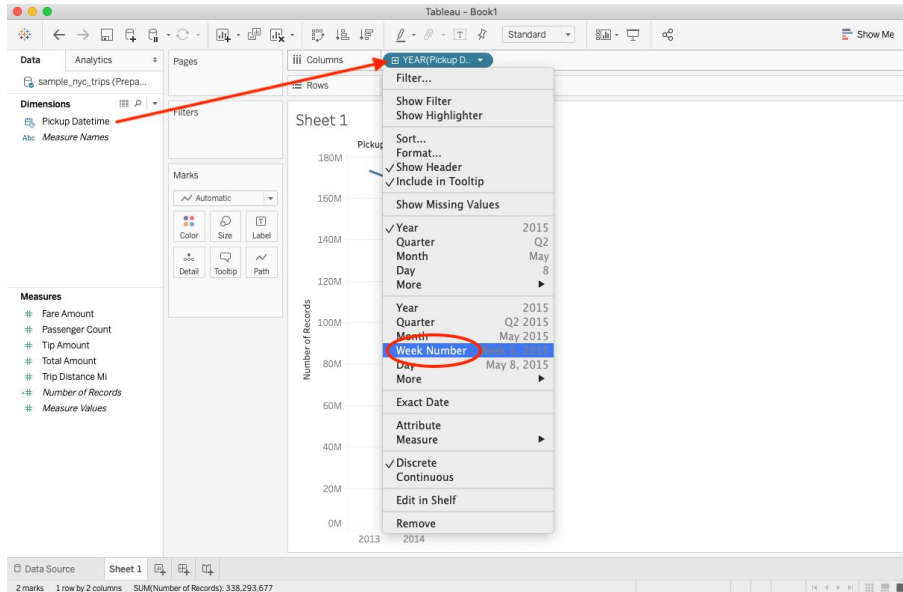
5. Create and publish a Tableau workbook using Tableau's "Impersonate using embedded password" configuration

When the Tableau workbook is opened, it will prompt you for your Dremio login credentials. Enter your username and password. The Tableau workbook will connect to Dremio and access the `Preparation.sample_nyc_trips` dataset and display the Dimensions and Measures associated with the dataset.

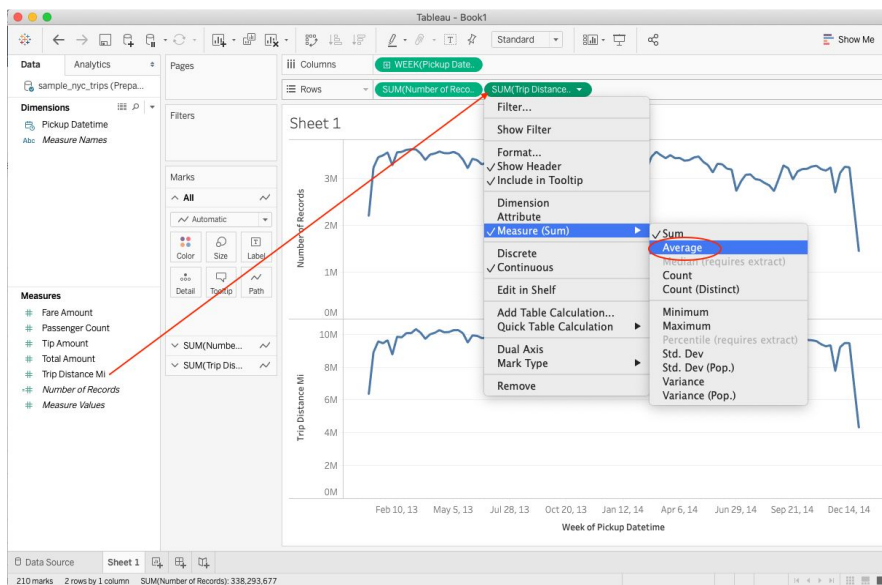
Add the number of taxi rides to the worksheet by dragging the `Number of Records` measure to the **Rows** section of the worksheet and you can see that there are over 1 billion records in this Dremio dataset.



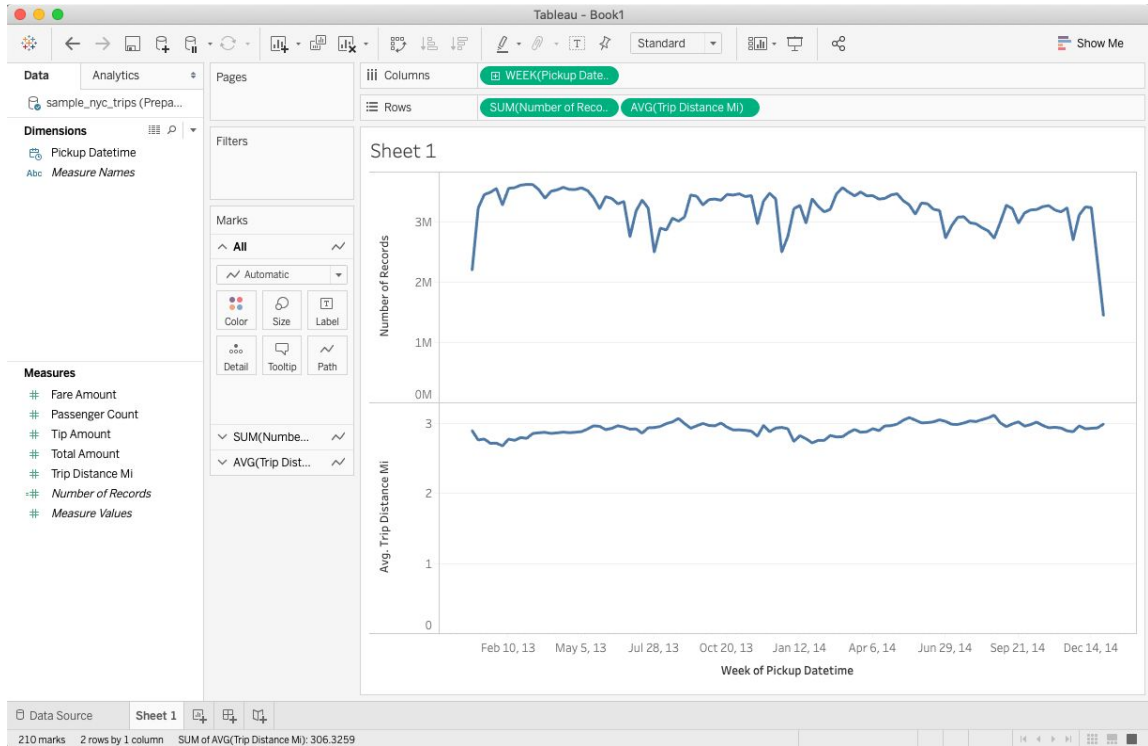
Next, add a graph of the number of taxi rides per week distributed across the years in the dataset. Drag the `Pickup Datetime` dimension to the **Columns** section of the worksheet, then select the **Week Number** group specification.



Now add a graph of the average trip distances by dragging the **Trip distance Mi** measure to the **Rows** section of the worksheet.




The final result will be a Tableau worksheet that plots the weekly taxi ride counts and the average trip distance of those taxi rides, from February 2013 through December 2014.



Now you are ready to publish the Tableau workbook into Tableau Server using the Tableau “Impersonate with embedded password” option.

First, sign into the Tableau Server. In the Tableau Desktop menu, select **Server > Sign in** and enter your Tableau Server address, then click on the **Connect** button.




Enter your Tableau Server username and password and click on the **Sign In** button.

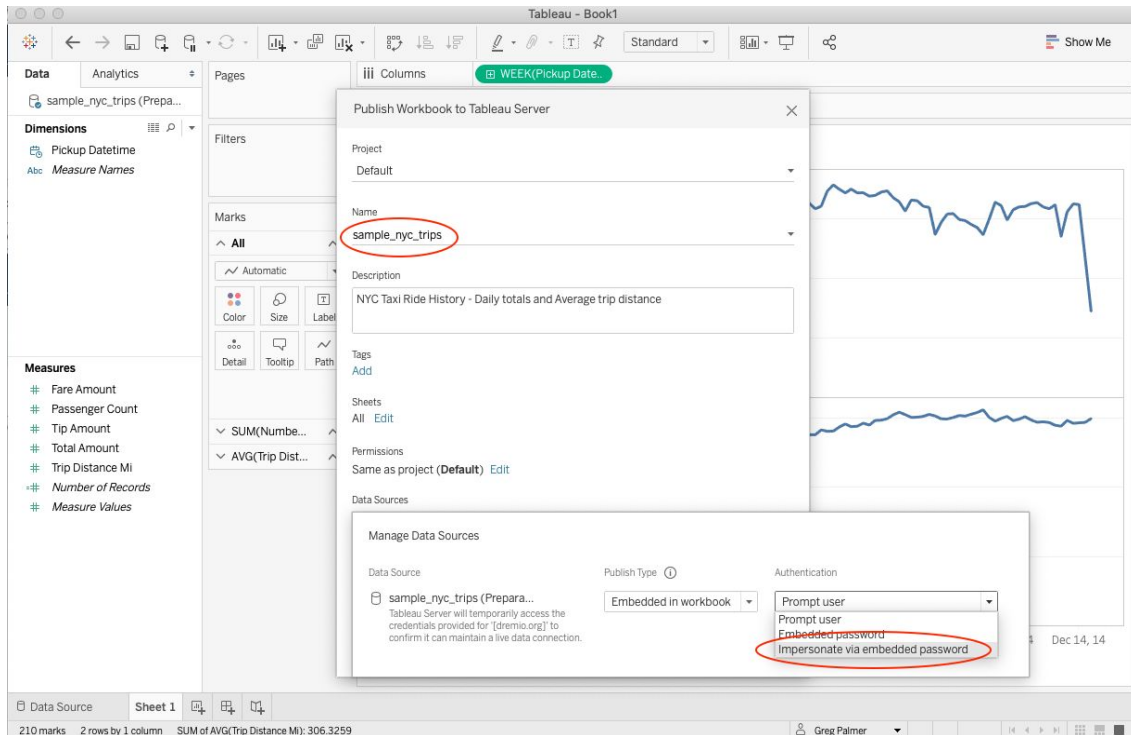


The image shows the Tableau login interface. At the top is the Tableau logo, which consists of a cluster of plus signs followed by the word "tableau" in a lowercase, sans-serif font. Below the logo are two input fields. The first is labeled "Username" and contains the text "greg@dremio.com". The second is labeled "Password" and contains a series of dots. Both fields have a small 'X' icon in the top right corner. Below the password field is a dark blue button with the text "Sign In" followed by a right-pointing arrow.

Now, publish your workbook by choosing the **Server > Publish Workbook** menu option from the Tableau Desktop menu.

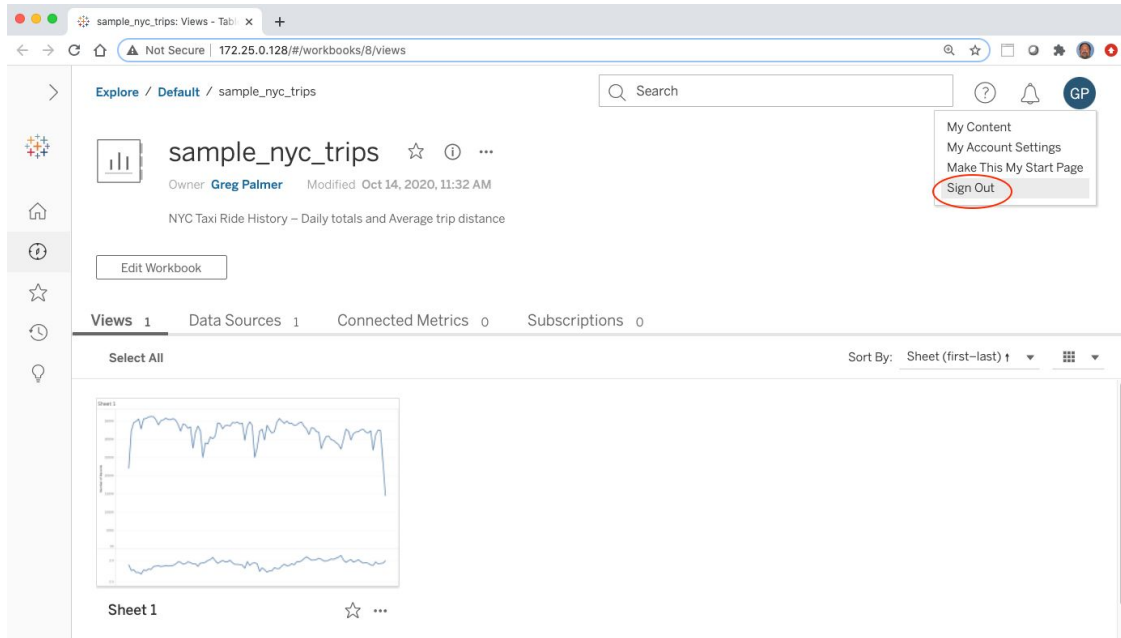


In the **Publish Workbook to Tableau Server** screen, enter a name for the workbook and click on the **Edit** link next to the data sources.



In the **Manage Data Sources** section, select the Authentication method **Impersonate via embedded password**. Then click on the **Publish** button. Tableau Desktop will open a Web browser window and display the published workbook.

Click the **Done** button and then sign out from the Tableau Server.



6. Run a Tableau workbook in Tableau Server as a different user using Dremio's “Inbound Impersonation” configuration

Now you will sign into Tableau Server as a different user from the one in which you published the workbook. Note that this user must also have access to Dremio via the Dremio inbound impersonation capability (see Step 3. above).

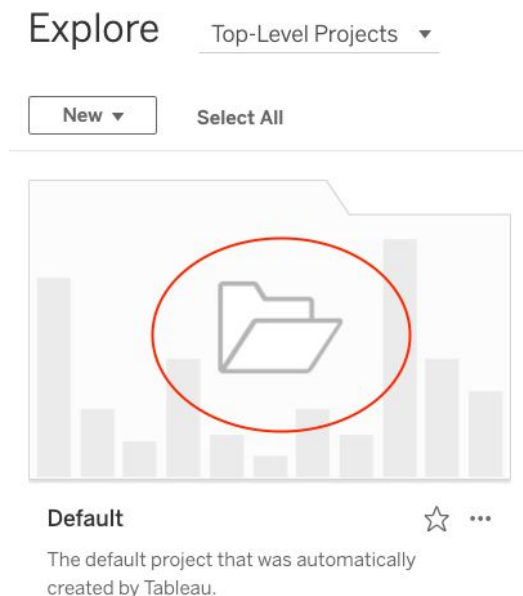
The image shows the Tableau login interface. At the top is the Tableau logo, which consists of a grid of plus signs followed by the word 'tableau' in a lowercase, sans-serif font. Below the logo are two input fields: 'Username' and 'Password'. The 'Username' field contains the text 'ben@dremio.com' and has a small 'X' icon to its right. The 'Password' field contains a series of dots and also has a small 'X' icon to its right. Below these fields is a dark blue button with the text 'Sign In' followed by a right-pointing arrow.

Once you have signed in to Tableau Server, click on the **Start Exploring** button on the Home page.

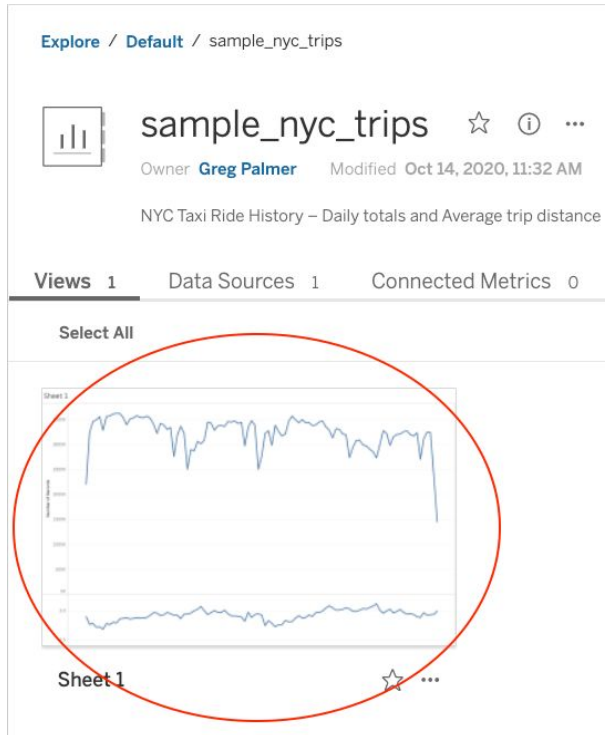
Home




Then open the project that you published your workbook in.



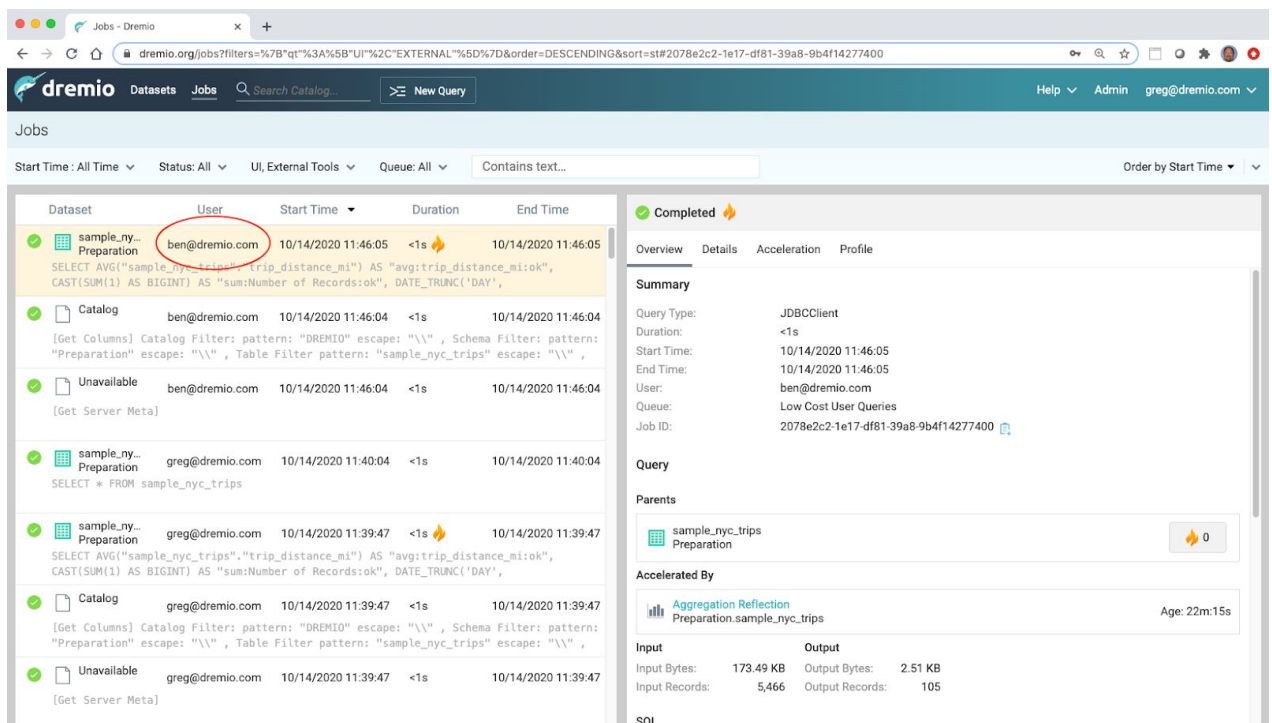
In the project folder, you will see the recently published workbook. To open the workbook, click on the `sample_nyc_trips` workbook and then double-click on the view for Sheet 1.



The workbook will be opened and you will see the same graphs that you saw when using Tableau Desktop, the number of taxi rides by week and the average trip distance.



Now, view the Dremio Jobs screen and notice that the query sent to Dremio by the Tableau Server based workbook was submitted as your current user and not the proxy user that published the workbook. And you were not prompted for a username or password to connect to Dremio. This is because Tableau already knows who you are and that the workbook data source was set up to authenticate with the “Impersonate via embedded password” option. Additionally, Dremio was configured to allow the proxy user (greg@dremio.com) to submit a query as the real user (ben@dremio.com).



The screenshot shows the Dremio Jobs screen in a web browser. The top navigation bar includes 'dremio', 'Datasets', 'Jobs', and a search bar. The 'Jobs' tab is active, displaying a table of jobs. The first job is highlighted, showing it was completed by user 'ben@dremio.com'. The right sidebar provides a detailed view of this job, including a summary, query, and acceleration details.

Dataset	User	Start Time	Duration	End Time
sample_ny... Preparation	ben@dremio.com	10/14/2020 11:46:05	<1s	10/14/2020 11:46:05
Catalog	ben@dremio.com	10/14/2020 11:46:04	<1s	10/14/2020 11:46:04
Unavailable	ben@dremio.com	10/14/2020 11:46:04	<1s	10/14/2020 11:46:04
sample_ny... Preparation	greg@dremio.com	10/14/2020 11:40:04	<1s	10/14/2020 11:40:04
sample_ny... Preparation	greg@dremio.com	10/14/2020 11:39:47	<1s	10/14/2020 11:39:47
Catalog	greg@dremio.com	10/14/2020 11:39:47	<1s	10/14/2020 11:39:47
Unavailable	greg@dremio.com	10/14/2020 11:39:47	<1s	10/14/2020 11:39:47

Completed Job Details:

- Summary:** Query Type: JDBCClient, Duration: <1s, Start Time: 10/14/2020 11:46:05, End Time: 10/14/2020 11:46:05, User: ben@dremio.com, Queue: Low Cost User Queries, Job ID: 2078e2c2-1e17-df81-39a8-9b4f14277400
- Query:** SELECT AVG("sample_nyc_trips"."trip_distance_mi") AS "avg:trip_distance_mi:ok", CAST(SUM(1) AS BIGINT) AS "sum:Number of Records:ok", DATE_TRUNC('DAY',
- Accelerated By:** Aggregation Reflection, Preparation.sample_nyc_trips, Age: 22m:15s
- Input/Output:** Input Bytes: 173.49 KB, Output Bytes: 2.51 KB, Input Records: 5,466, Output Records: 105

Summary

This document discussed why you would use the Dremio Connector for Tableau and how to set it up on the desktop as well as on the server. It also walked you through the steps needed to create a Tableau workbook that utilizes the connector.

For further reading and for reference, here are some useful links:

<http://docs.dremio.com/client-applications/tableau-sdk.html>

<https://extensiongallery.tableau.com/connectors>

<https://tableau.github.io/connector-plugin-sdk/#>